Building a Chinese Wall in Security-Enhanced Linux

by

Calum Cardownie

In partial fulfilment of the
requirements for the degree of
Bachelor of Science
in
CS408: Individual Project

University of Strathclyde Glasgow

Department of
Computer and Information Sciences

March, 2026

Abstract

The abstract is a bird's eye view of the project.

It should not exceed one page. Mention the scope and objectives of the project, the methodology used,the main findings, and significance of your results.

Acknowledgements

I would like to thank

Contents

1 Introduction **1**
1.1 Aims . . . 1
1.2 Objectives . . . 2
1.3 Report Format . . . 2

2 Background Knowledge **3**
2.1 The Chinese Wall (Brewer-Nash) Origin . . . 3
2.2 The Chinese Wall (Brewer-Nash) Intricacies . . . 4
2.3 Security-Enhanced Linux . . . 4
2.4 Bell-LaPadula model in SELinux . . . 5
2.5 Chinese Wall in UNIX . . . 5
2.6 Hypervisor-Level Chinese Wall . . . 5
2.7 Chinese Wall in Worlflow Management . . . 6

3 Specification **7**
3.1 Methodology . . . 7
3.2 Analysis . . . 8
3.3 Requirements . . . 8

4 Design **11**
4.1 Environment Design . . . 11
4.2 System Design . . . 11

5 Product **13**
5.1 Implementation . . . 13

6 Verification & Validation **14**
6.1 Testing Suite . . . 14
6.2 Testing Constraints . . . 14

7 Results & Evaluation **15**
7.1 Evaluation Process . . . 15
7.2 Results of Evaluation . . . 15
7.3 Returning to the Research Questions . . . 15

8 Discussion & Reflection **16**
8.1 Interpreting the Results . . . 16
8.2 Reflection . . . 16
8.3 Challenges . . . 16
8.4 Limitations . . . 16
8.5 Future Work . . . 16

9 Conclusion **17**

A Appendix **19**
A.1 Ethical Approval Form . . . 19

A.2 Participant Information Sheet 19
A.3 Consent Form 19
A.4 Marking Scheme 20

1 Introduction

1.1 Aims

The primary aim of this project is implementation of the Chinese Wall (Brewer-Nash) Model, by leveraging the mandatory access control system introduced by Security-Enhanced Linux. SELinux is a well-known solution to security concerns threatening servers and personal devices. SELinux implementing Red Hat (IBM) have solidified themselves as the industry choice for secure hosting solutions, being a leading operating system used by over 90% of Fortune 500 companies [1] for its quality privacy guarantee.

The Chinese Wall's dynamic nature makes it an intriguing security policy, which let it stand in stark contrast from typical static security policies. For this reason or others, it has received a noticeable lack of attention in the space surrounding SELinux. While providing an implementation for SELinux, this report delves into the use cases of this specific policy, where it has succeeded in the past, and what its inclusion in SELinux allows.

"The more you know, the less you can share" is the basis for this policy, where a subject (a process, daemon or user) in the system loses more privileges as it completes its lifecycle. Conflict of Interest (`CoI`) classes dictate the main restrictions, preventing the subject from sharing sensitive information between objects (files) that are conflicting with each-other. A lattice based system is used to transition subject labels based on their read and write history accordingly, which should suffice to handle all cases of the Chinese Wall.

Previous research into both Chinese Wall effectiveness and SELinux was factored in the design process, along with hands on experimentation and prototyping of different solutions. The development of the Chinese Wall daemon should serve to open possibilities of dynamic policies in the otherwise static SELinux, and potentially lead the way in future security solutions.

To enforce this policy model in SELinux, policy modules must be compiled into the kernel via the security server and work in tandem with the designed Chinese Wall Security Daemon, developed to be capable of enforce dynamic label alterations to monitored subjects in a directory. Finally, this allows the SELinux Active Vector Cache (AVC) to handle the restrictions on a kernel level, successfully interfering in any attempted `CoI` breaches.

1.2 Objectives

1.1 Create a suitable virtual environment for SELinux.

1.2 Understand the principles of:

1.2.A The Chinese Wall security model.

1.2.B SELinux's security server & type enforcement system.

1.3 Implement a static implementation of the policy, confined by SELinux.

1.4 Develop a system daemon capable of monitoring subjects and their operations on objects.

1.5 Update said daemon to dynamically alter labels, allowing SELinux to enforce the Chinese Wall policy.

1.6 Fully integrate the Chinese Wall Daemon into SELinux's environment, including its own policy module, file hierarchy compliance and good UNIX citizenship.

1.7 Analyse the difficulties that arose with development and conclude why development in this area has been lacklustre.

1.3 Report Format

The report is structured in an sequence appropriate to the process undertook addressing this challenge, detailing background research, solution design, implementation and testing, followed by results and evaluation.

2 Background Knowledge

2.1 The Chinese Wall (Brewer-Nash) Origin

The Chinese Wall security model - or Brewer and Nash model [3] - is a security model that prevents information from conflicts of interest (`CoI`) groups being spread. As opposed to the Bell LaPadula (BLP) model, which enforces "no read up, no write down" [2], the Chinese Wall model allows for information flow freely until a subject (e.g. a user or process) attempts to read from a `CoI` object (piece of information). Upon this access, the system dynamically revokes their write permissions to other objects in the `CoI` group, creating a "wall" in-between them. Consider the following scenario:

A data server stores information on Banks A and B, and Oil Company A and B. These are in `CoI` groups with their respective counterpart. If a user approaches the dataset, they are initially able to access any object. As far as the system knows, this user is free as they have no prior knowledge of either companies. Suppose the user has read from Bank A, the situation develops - the user now possesses information that, deemed by the Chinese Wall, should not be allowed to reach Bank B, as they are conflicting with Bank A. This is where the wall revokes the write access from the user in regard to Bank B. Now information *cannot* flow from Bank A to B.

- **BN Read Rule**: Subject `S` can read object `O` only if `O` is in the same Company Dataset (`CD`) as an object already read by `S`, or if `O's` (`CD`) is in a different Conflict of Interest (`CoI`) group from all `CD`s of objects already read by `S`.
- **BN Write Rule**: Subject `S` can write to an object `O` only if `S` can read `O` **and** all objects `O'` read by `S` are in the same `CD` as `O`.

The Brewer Nash model provides a mathematical format that is based on the larger Chinese Wall policy, a pillar of compliance in the financial sector. Chinese Walls are a common practice that should theoretically and practically prevent misuse of probable information leakage in banks [6] and other financial management firms. After the Brewer Nash model was communicated, another paper was published once again by Brewer and Nash, but also TY Lin, who presented them with potential improvements on their model [8], leading to the Aggressive Chinese Wall. This idea, trying to disprove the assumption that the conflict of interest is an equivalent relation, has been accepted into the mainstream since it was proposed. This is the type of implementation the 1997 UNIX proof believed in [5], illustrating the possibility of its creation in UNIX. Such proofs in UNIX lead it to be

assumed to theoretically be possible in Linux and SELinux. This proof indicates there is research, but there is a lack of documented attempts to specifically implement it into the UNIX or adjacent kernel.

2.2 The Chinese Wall (Brewer-Nash) Intricacies

2.3 Security-Enhanced Linux

Security Enhanced Linux implements a Mandatory Access Control (MAC) system, providing a policy-driven security layer extended beyond the classic Discretionary Access Control (DAC) system that UNIX based systems use. DAC is defined such that objects and their privileges are managed by whoever made them, and the root user can access and control anything. DAC in UNIX leverages the `chmod` (change mode) command, allowing read, write and execute access of objects to be edited for the 3 main groups subjects can belong to. While DAC brings about great flexibility, the models' reliance on ownership often leads to uncontrolled privilege propagation. It is known to be difficult to measure the security of the DAC model, and in some cases it is undecidable [4]. MAC is the other model that shares the spotlight with DAC in system security. It utilises, labels, attached to both subjects and objects. Installed policies into the security server govern what labels can interact. Labels are devised in the form `user_u:role_r:type_t:c0` For instance, a base policy may be that an object of type `httpd_t` can only be operated on by a subject of type `httpd_t` and not `sshd_t`. This provides strong mitigation against broken access control damages, as subjects can only access objects that they should and are unable to change the policy controls. For this reason SELinux is adopted by many distributions of Linux, including Red Hat, Fedora, CentOS and even mobile operating system giant Android. SELinux is based off of the Flux Advanced Security Kernel (FLASK) security architecture for flexible nondiscretionary access controls, where the architecture is split into policy enforcement and policy decision-making. SELinux's origin can be attributed to the National Security Agency's Trusted UNIX (TRUSIX) Working Group, working together to publish a Rainbow Book (A series of Computer Security standards and guidelines by the United States Government), and created a second (unpublished) book that included a formal SELinux model and associated evaluation evidence prototype. SELinux was a demonstration that mandatory access control was viable and could be added to the Linux kernel, and so it was merged into the mainline kernel in the 2.6 series of the kernel. Linux was able to adopt this because of its derivation from UNIX, Linux being an open source alternative to the UNIX that was locked down by ATT. Linus

Torvalds combined his Linux kernel with the GNU system tools to create GNU/Linux, leading to what we now know as Linux. This open-source philosophy and flexibility were key factors enabling SELinux's integration into the mainline kernel in the 2.6 series.

2.4 Bell-LaPadula model in SELinux

2.5 Chinese Wall in UNIX

2.6 Hypervisor-Level Chinese Wall

Cloud computing is an avenue that has had a successful Chinese Wall model implementation [9]. In this case, there are many Virtual Machines (VMs) running on one system, with the virtual machine monitor, or "hypervisor" acting as the delegating force of the computers' hardware. Each virtual machine is a separate operating system, commonly being completely independent of the other VMs on a machine. Inter-VM attacks can happen if a VM on the system is breached, and can lead to one infected VM spreading to the whole system. This is where the Chinese Wall was created, to eliminate the possibility of inter-VM attacks from competitors. The security model was implemented on the Kernel-base Virtual Machine (KVM) level. Once a VM in a `CoI` group is running on a physical host, the Wall prevents other VMs belonging to the same `CoI` group from being scheduled on the same machine. This effectively eliminates the risk of inter-VM attacks in these systems.

This section should provide the grounding for your project. You will be able to refer back to this in Section 3, where you report on design decisions and methodology.

There should be a strong link to your research questions (for project types 2 and 3). There will be a strong link to the software design decisions for project type 1.

Don't just include a paragraph for each paper you read. Synthesize it and create a story line.

Review between 5 and 10 authoritative sources. For research based projects, some of these should be from peer-reviewed journals or conferences.

This section will reference literature like this [**?**]. Sloppy referencing will (1) take time and effort, and (2) lose you marks if you don't do it.

2.7 Chinese Wall in Worlflow Management

Beyond virtualized environments, the Chinese Wall model has also successfully been implemented in workflow management systems. Workflow Management Systems (WfMS) are intrinsically network-based applications, potentially hosting multiple conflicting sources of information. As such, a secure WfMS must implement security features, specifically targeting information confidentiality as a main point of effort. The Chinese Wall security model can be leveraged to prevent conflict of interest leakage. A Chinese Wall was implemented using role-based access control (RBAC) to enforce, and Java to handle the Chinese Wall logic [7]. The paper mentions that RBAC is unable to implement general cases of the Chinese Wall, as it has issues implementing the BN read rule in dynamic contexts.

3 Specification

3.1 Methodology

3.1.1 Agile

A standard agile methodology was followed while completing this project. Agile methodology focuses on multiple iterative stages, constantly evaluating the progress and direction of the project, and tuning them to reach the solution. As SELinux is a system with predefined capabilities and limitations, various different methods were used throughout to understand the boundaries and prototyping was updated accordingly. Experimentation in SELinux was carried out simultaneously alongside research, allowing accumulated insight to be backed by personal experience, and a more complete and satisfying solution.

This agile process was accompanied by a simple kanban board, easily conveying the stage of the project as well as next steps and currently occupied tasks. This allowed for a more streamlined lifecycle of an idea.

A Gantt Chart was constructed in advance to demonstrate the intended project timeline:

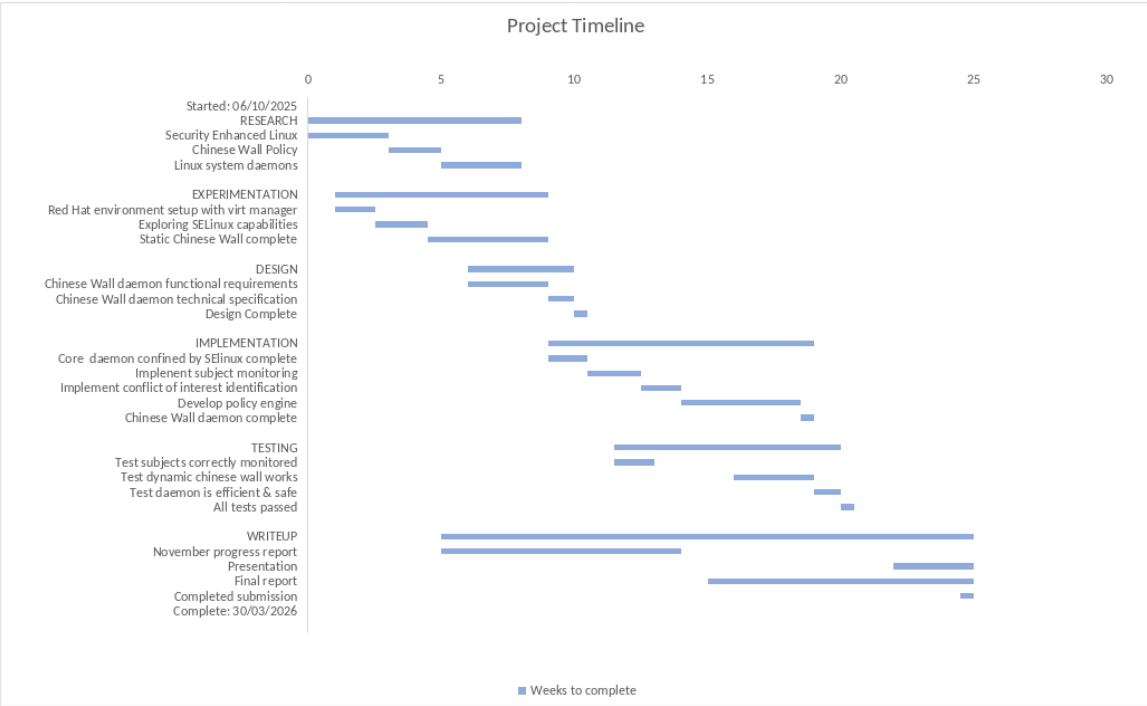

Figure 1: Project Gantt chart

3.1.2 Issues with Waterfall

Agile stands in contrast to the alternative methodology - waterfall - in which involves a predefined set of requirements solely working towards the final objective. This methodology benefits experience, where familiarity with the project environment pays off. It relies on decisions to be made in advance with confidence, of which more experience with the environment would be needed. As this is a project within time constraints and a single developer, agile was selected.

3.2 Analysis

To enforce a dynamic security policy, its natural to deduce that some sort of process must be constantly running to be capable of enforcing said dynamism. During experimentation on the boundaries of SELinux, it was surmised that the most effective way to leverage its MAC system was to change the type fragment of the labels on both subjects and objects. The user, role and category of the subjects/objects were not sufficient to enforce a security policy, as policy modules were unable to be compiled to enforce restrictions based entirely on said label fragments.

Aggressive experimentation on the type fragment of the label opened up the philosophy that would be the main premise of the Chinese Wall Daemon. Pre-compiled policy modules for all scenarios, and a daemon that altered the labels of subjects at runtime according to their behaviour and the Chinese Wall policy. This daemon would also need to be fully integrated into SELinux's system, running under its own `cwalld_t` type rather than `unconfined_t`.

Policy modules are unable to be added to the SELinux security server at runtime, as it requires a complete re-compile of the security server and kernel. This operation is capable of taking anywhere between 4 and 8 seconds - suffice to say unrealistic at runtime for something intended as a background process.

3.3 Requirements

Requirements specified early in the agile cycle follow the MoSCoW principle. Requirements are split into:

- Must have - features that must be implemented to achieve the goal.
- Should have - features that should be implemented to develop the solution.

- Could have - features that could be implemented if there was more time.
- Won't have - features that were considered but decided against.

Functional requirements:

The Chinese Wall Daemon (`cwalld`) Must Have:

- `cwalld` **must** alter necessary configurations in the system before running.
- `cwalld` **must** monitor all operations made by subjects across a given directory.
- `cwalld` **must** make decisions based on subject label and object label aligned with the Chinese Wall policy.
- `cwalld` **must** alter the labels of subjects at upon call.
- `cwalld` **must** be accompanied by policy modules for objects.
- `cwalld` **must** be accompanied by policy modules for subjects.
- `cwalld` **should** be confined by SELinux with its own custom policy.
- `cwalld` **should** run as a service in the system with its own service declaration file.
- `cwalld` **should** constantly write operations to a log.
- `cwalld` **should** include a binary to follow and display the log in real time.
- `cwalld` **could** automatically create necessary policy modules for its runtime.
- `cwalld` **could** automatically install necessary policy modules before running
- `cwalld` **could** dynamically enforce the policy given a set of `CoI` classes as input.
- `cwalld` **won't** enforce the Chinese Wall policy on the system-wide.

Non-functional requirements:

- `cwalld` **shall** run as a daemon in user-space of the computer.
- `cwalld` **shall** solely be controlled by the super-user.
- `cwalld` **shall** be written with structural, intentional and idiomatic code.
- `cwalld` **shall** be a good UNIX citizen, following classic system conventions.

- `cwalld` **shall** function as intended across both x86 and ARM architecture.

These user stories were decided solidified after considering the functionality that would be needed to enforce the Chinese Wall. While the user stories were decided from the beginning, the means of which to complete them was left open, in line with agile structure.

4 Design

Describe all details of the design and procedures used to achieve the project objectives. Do this chronologically.

4.1 Environment Design

The SELinux AVC makes

4.2 System Design

Show how you designed your database (if appropriate) and how you designed your system architecture, and the individual parts. Use UML and an Entity Relationship diagram

It should be detailed enough to allow for an assessment of the rigour of your process, and, in the case of research projects, in terms of how well grounded your research is in the research literature. In these cases, refer back to relevant sections in the previous chapter.

Say which software lifecycle approach you used e.g., Waterfall, Spiral, Agile.

How did you gather user requirements?

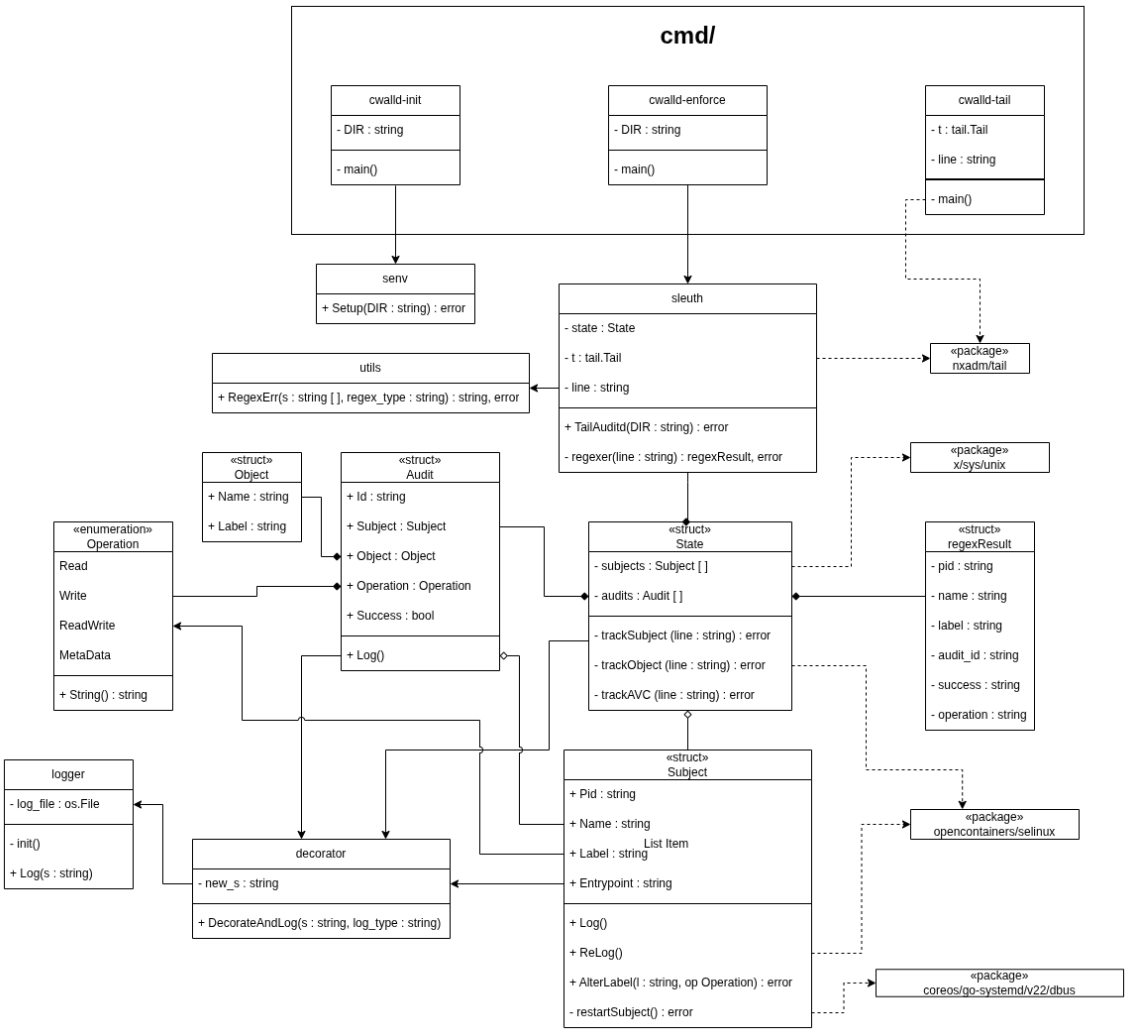

Figure 2: cwalld UML diagram

5 Product

5.1 Implementation

Provide implementation details: language used, architecture (e.g. server and client, or hub and spoke). Explain how you secured personal details.

6 Verification & Validation

6.1 Testing Suite

6.2 Testing Constraints

How did you verify that your software was debugged and worked correctly?

How did you validate that you had implemented all the functional and non-functional requirements?

7 Results & Evaluation

Results are often presented in tables, figures and other relevant illustrations. Include text that refers to these figures/tables.

7.1 Evaluation Process

If you involved humans in the evaluation, how many did you have? What can you say about the demographics of your participants? (If you did collect these)

In terms of the *user interface*, how did you carry out a usability evaluation, how did you go about doing this? How did you recruit participants?

In terms of delivering *functionality*, did you carry out user acceptance testing? (see attached guidance).

7.2 Results of Evaluation

This Section includes a direct interpretation of the gathered data and evaluation processes.

7.3 Returning to the Research Questions

Return to research questions or objectives as appropriate.

7.3.1 RQ1

It is clear from our findings that James Bond was born in Wigtown in Scotland. However, he grew up in Diss, in Norfolk. We know this because ....

7.3.2 RQ2

We were not able to answer this question from our studies, although some suggestions were made. These could not be proven.

7.3.3 Objective 1

7.3.4 Objective 2

8 Discussion & Reflection

8.1 Interpreting the Results

Here you will discuss your findings. This is especially relevant for research projects. You might interpret what the data and evaluation implies, both for future research and for practice (if appropriate).

The discussion is **not** a review of literature. You should try to compare research findings with previous work, provide explanations for your findings, discuss research findings, in terms of their contribution.

8.2 Reflection

Look back and think about what you would do differently if you were going to start the project with the knowledge you have now. Be honest about your mistakes or missteps.

8.3 Challenges

This is not the place to mention personal circumstances but rather challenges related to the work involved in the project.

8.4 Limitations

Acknowledge things like: small number of participants, software wasn't completely debugged, or whatever else went wrong and affected your project. *Include as appropriate*

8.5 Future Work

If someone else wanted to build on your project's product, what would be cool to do next?

9 Conclusion

The conclusion is similar to when a plane lands. You don't rewrite the introduction. You say something like - I addressed the problem outlined in the introduction, and I built some software to do this. I tested the software like this ADD FEW WORDS.

Summarize main findings drawn from the project work. Mention the objectives or research questions. Do not repeat points raised in the discussion and reflection Section. If applicable, you can make recommendations. The conclusion should NOT contain any references.

References

[1] Red hat enterprise linux | red hat partner connect.

[2] D. E. Bell and L. J. Lapadula. Secure computer systems: Mathematical foundations. Number: MTR2547VOL1.

[3] D. Brewer and M. Nash. The chinese wall security policy. In *Proceedings. 1989 IEEE Symposium on Security and Privacy*, pages 206–214. IEEE Comput. Soc. Press.

[4] Y. Fan, Z. Han, J. Liu, and Y. Zhao. A mandatory access control model with enhanced flexibility. In *2009 International Conference on Multimedia Information Networking and Security*, volume 1, pages 120–124. ISSN: 2162-8998.

[5] S. N. Foley. Building chinese walls in standard unixTM. 16(6):551–563.

[6] L. Herzel and D. E. Colling. The chinese wall and conflict of interest in banks. 34(1):73–115. Publisher: American Bar Association.

[7] Y.-C. Hsiao and G.-H. Hwang. Implementing the chinese wall security model in workflow management systems. In *International Symposium on Parallel and Distributed Processing with Applications*, pages 574–581. ISSN: 2158-9208.

[8] T. Y. Lin. Chinese wall security policy-an aggressive model. pages 282,283,284,285,286,287,288,289–282,283,284,285,286,287,288,289. IEEE Computer Society.

[9] T.-H. Tsai, Y.-C. Chen, H.-C. Huang, P.-M. Huang, and K.-S. Chou. A practical chinese wall security model in cloud computing. In *2011 13th Asia-Pacific Network Operations and Management Symposium*, pages 1–4.

A Appendix

This is where you can include your documentation.

Remember that the marker is not required to read this, but might well check to ensure that you have included product documentation, and ethical approval, as required.

A.1 Ethical Approval Form

If your project required you to do any evaluation with humans, you MUST include this. It can be downloaded from the Ethics system.

`https://local.cis.strath.ac.uk/wp/extras/ethics/index.php`

A.2 Participant Information Sheet

If your project required you to do any evaluation with humans, you MUST include this

`https://www.strath.ac.uk/ethics/informationsheetandconsentform/`

A.3 Consent Form

If your project required you to do any evaluation with humans, you MUST include this.

`https://www.strath.ac.uk/ethics/informationsheetandconsentform/`

A.4 Marking Scheme

REMEMBER TO DELETE THIS. IT IS ONLY INCLUDED FOR your INFORMATION.

11.2. Marking Schemes

There are three marking schemes which weight the assessment criteria as follows.

	Software Development Based	**Experimentation-based with Significant Software Development**	**Experiment-based**
Project progress presentation	10%	10%	10%
Product			
Implementation (including documentation as this indicates maintainability)	25	20	10
Verification and Validation	10	5	5
Product Total:	35%	25%	15%
Process			
Methodology, analysis and documentation	15	20	20
Design	10	5	5
Process Total	25%	25%	25%
Results and Evaluation			
	15%	25%	35%
Report Presentation			
	10%	10%	10%
Student Performance			
	5%	5%	5%